

#Importing necessary python library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

#Uploading Automobile Dataset

```
from google.colab import files
uploaded = files.upload()

<IPython.core.display.HTML object>

Saving Automobile_data.csv to Automobile_data.csv
```

#Load the dataset

```
df = pd.read_csv('Automobile_data.csv')
```

#Univariate Analysis Numeric Variables Preparing Data

```
df = df.replace('?', pd.NA)

# convert data types
df['normalized-losses'] = pd.to_numeric(df['normalized-losses'],
errors='coerce')
df['bore'] = pd.to_numeric(df['bore'], errors='coerce')
df['stroke'] = pd.to_numeric(df['stroke'], errors='coerce')
df['horsepower'] = pd.to_numeric(df['horsepower'], errors='coerce')
df['peak-rpm'] = pd.to_numeric(df['peak-rpm'], errors='coerce')
df['price'] = pd.to_numeric(df['price'], errors='coerce')

# drop rows with missing values
df = df.dropna()

# reset index
df = df.reset_index(drop=True)

# display the first few rows
print(df.head())
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	2	164.0	audi	gas	std		

```

four
1      2      164.0 audi      gas      std
four
2      1      158.0 audi      gas      std
four
3      1      158.0 audi      gas      turbo
four
4      2      192.0 bmw      gas      std
two

    body-style drive-wheels engine-location wheel-base ... engine-
size \
0      sedan      fwd      front      99.8 ...
109
1      sedan      4wd      front      99.4 ...
136
2      sedan      fwd      front      105.8 ...
136
3      sedan      fwd      front      105.8 ...
131
4      sedan      rwd      front      101.2 ...
108

    fuel-system bore stroke compression-ratio horsepower peak-rpm
city-mpg \
0      mpfi  3.19   3.4      10.0      102.0      5500.0
24
1      mpfi  3.19   3.4      8.0      115.0      5500.0
18
2      mpfi  3.19   3.4      8.5      110.0      5500.0
19
3      mpfi  3.13   3.4      8.3      140.0      5500.0
17
4      mpfi  3.50   2.8      8.8      101.0      5800.0
23

    highway-mpg price
0      30  13950.0
1      22  17450.0
2      25  17710.0
3      20  23875.0
4      29  16430.0

[5 rows x 26 columns]

df.describe()

    symboling normalized-losses wheel-base length
width \
count  159.000000      159.000000  159.000000  159.000000

```

```

159.000000
mean    0.735849      121.132075    98.264151    172.413836
65.607547
std     1.193086      35.651285     5.167416     11.523177
1.947883
min     -2.000000      65.000000     86.600000    141.100000
60.300000
25%     0.000000      94.000000     94.500000    165.650000
64.000000
50%     1.000000      113.000000    96.900000    172.400000
65.400000
75%     2.000000      148.000000    100.800000   177.800000
66.500000
max      3.000000      256.000000    115.600000   202.600000
71.700000

```

```

          height  curb-weight  engine-size      bore      stroke  \
count  159.000000   159.000000   159.000000   159.000000   159.000000
mean    53.899371  2461.138365   119.226415    3.300126    3.236352
std      2.268761   481.941321    30.460791    0.267336    0.294888
min     49.400000  1488.000000    61.000000    2.540000    2.070000
25%     52.250000  2065.500000    97.000000    3.050000    3.105000
50%     54.100000  2340.000000   110.000000    3.270000    3.270000
75%     55.500000  2809.500000   135.000000    3.560000    3.410000
max     59.800000  4066.000000   258.000000    3.940000    4.170000

```

```

          compression-ratio  horsepower      peak-rpm    city-mpg
highway-mpg  \
count      159.000000   159.000000   159.000000   159.000000
159.000000
mean         10.161132   95.836478  5113.836478   26.522013
32.081761
std          3.889475   30.718583   465.754864    6.097142
6.459189
min          7.000000   48.000000  4150.000000   15.000000
18.000000
25%          8.700000   69.000000  4800.000000   23.000000
28.000000
50%          9.000000   88.000000  5200.000000   26.000000
32.000000
75%          9.400000  114.000000  5500.000000   31.000000
37.000000
max         23.000000  200.000000  6600.000000   49.000000
54.000000

```

```


          price
count    159.000000
mean   11445.729560
std    5877.856195
min    5118.000000

```

25%	7372.000000
50%	9233.000000
75%	14719.500000
max	35056.000000


#Categorical Variables

```
sns.histplot(x='length', data=df)  
<Axes: xlabel='length', ylabel='Count'>
```



#Bi-Variate Analysis Numerical Variables


```
sns.displot(df[['price', 'width', 'length', 'curb-weight']])  
plt.show()
```



#Categorical Variables

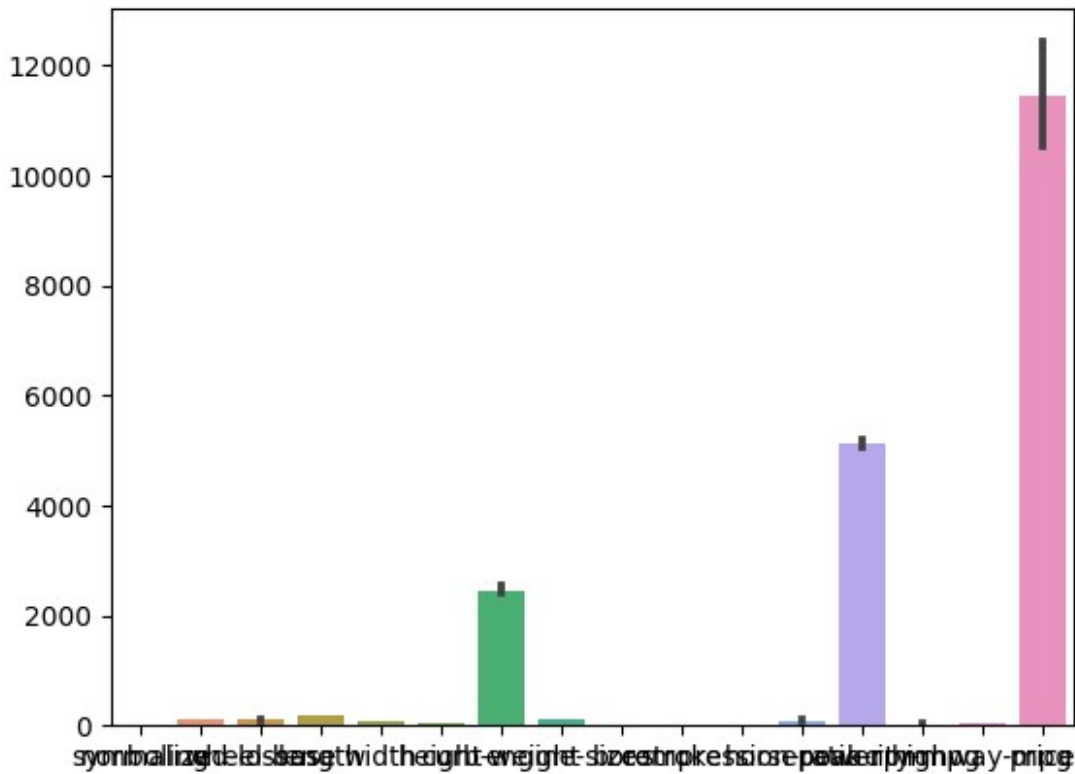
```
plt.figure(figsize=(10, 8))
sns.displot(x='make', y='num-of-cylinders', data=df)
plt.xticks(rotation=90)
plt.title('num-of-cylinders')
plt.show()
```

<Figure size 1000x800 with 0 Axes>



#Muliti-Variate Analysis

```
sns.barplot(df)
plt.show()
```



#Task 4: Perform data preprocessing handling missing values

```
df.replace('?', np.nan, inplace=True)
df.dropna(inplace=True)
```

#Handling categorical values

```
le = LabelEncoder()
df['make'] = le.fit_transform(df['make'])
df['fuel-type'] = le.fit_transform(df['fuel-type'])
df['aspiration'] = le.fit_transform(df['aspiration'])
df['num-of-doors'] = le.fit_transform(df['num-of-doors'])
df['body-style'] = le.fit_transform(df['body-style'])
df['drive-wheels'] = le.fit_transform(df['drive-wheels'])
df['engine-location'] = le.fit_transform(df['engine-location'])
df['engine-type'] = le.fit_transform(df['engine-type'])
df['num-of-cylinders'] = le.fit_transform(df['num-of-cylinders'])
df['fuel-system'] = le.fit_transform(df['fuel-system'])
```

#perform scaling

```
scaler = StandardScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df.drop('price',
axis=1)), columns=df.columns[:-1])
df_scaled['price'] = df['price']
```

#Task 5: Build Machine Learning Model

```
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
import numpy as np

# Load the dataset
df = pd.read_csv("Automobile_data.csv")

# Handling missing values
df.replace('?', np.nan, inplace=True) # replace '?' with NaN
df['normalized-losses'].fillna(df['normalized-losses'].median(),
inplace=True) # replace missing values with median
df['num-of-doors'].fillna(df['num-of-doors'].mode()[0], inplace=True)
# replace missing values with mode

# Encode categorical variables
df = pd.get_dummies(df, columns=['make', 'fuel-type', 'aspiration',
'num-of-doors', 'body-style',
                                'drive-wheels', 'engine-location',
'engine-type', 'num-of-cylinders',
                                'fuel-system'])

# Handle missing values
imputer = SimpleImputer(strategy='mean')
df = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)

# Perform scaling
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# Convert numpy array to DataFrame
df_scaled = pd.DataFrame(df_scaled, columns=df.columns)

# Split data into train and test sets
X = df_scaled.drop('price', axis=1)
y = df_scaled['price']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Build and fit the linear regression model
lr = LinearRegression()
lr.fit(X_train, y_train)

LinearRegression()
```

#Task 6: Evaluating Machine Learning Model


```
y_pred = lr.predict(X_test)
print('R-squared score:', r2_score(y_test, y_pred))
```

R-squared score: -1.7804692480809753e+23